

EE123 Course Notes

Anmol Parande

Spring 2020 - Professor Miki Lustig

Disclaimer: These notes reflect EE123 when I took the course (Spring 2020). They may not accurately reflect current course content, so use at your own risk. If you find any typos, errors, etc, please raise an issue on the [GitHub repository](#).

Contents

1	The DFT	3
1.1	Convolution and the DFT	4
1.1.1	Circular Convolution	4
1.1.2	Linear Convolution with the DFT	5
1.1.3	Block Convolutions	5
1.2	FFT	6
1.2.1	Decimation in Time	7
1.2.2	Decimation in Frequency	8
2	Spectral Analysis	8
2.1	Short Time Fourier Transform (STFT)	9
2.2	Discrete STFT	10
2.3	Time-Frequency Uncertainty	10
2.4	Wavelets	12
2.4.1	Discrete Wavelet Transform	13

3	Sampling	13
3.1	Ideal Sampling	13
3.1.1	Nyquist Theorem	15
3.1.2	Discrete Time Processing of a Continuous Time Signal	16
3.1.3	Continuous Time Processing of Discrete Time Signals	17
3.1.4	Downsampling	17
3.1.5	Upsampling	18
3.2	Multi-Rate Signal Processing	18
3.2.1	Exchanging Filter Order During Resampling	19
3.2.2	Polyphase Decomposition	19
3.3	Practical Sampling (ADC)	20
3.3.1	Quantization	21
3.4	Practical Reconstruction (DAC)	22
4	Filtering	23
4.1	Transform Analysis of LTI Systems	23
4.2	All Pass Systems	24
4.3	Minimum Phase Systems	24
4.4	Generalized Linear Phase Systems	25
4.5	Filter Design	26
4.5.1	Windowing Method	26
4.5.2	Optimal Filter Design	27

1 The DFT

Whereas the CTFT takes a continuous signal and outputs a continuous frequency spectrum and the DTFT takes a discrete signal and outputs a continuous, periodic frequency spectrum, the Discrete Fourier Transform takes a discrete finite signal and outputs a discrete frequency spectrum. This is useful for signal processing because we cannot store infinite signals in a computer's memory.

Definition 1 For a length N finite sequence $\{x[n]\}_0^{N-1}$, the Discrete Fourier Transform of the signal is a length N finite sequence $\{X[k]\}_0^{N-1}$ where

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-j \frac{2\pi}{N} kn}$$

One way to interpret the DFT is in terms of the Fourier series for a discrete periodic signal $\tilde{x}[n] = x[(n)_N]$ where the $((n))_N = n \bmod N$. Recall that the coefficient of the k th term of the Fourier Series is

$$a_k = \frac{1}{N} \sum_{n=0}^{N-1} x[n] e^{-j \frac{2\pi}{N} kn}$$

Notice that the a_k of the Fourier Series are the DFT values except scaled by a factor of N . In other words, if we extend a finite signal periodically, then the DFT and the DTFS are the same up to a constant scale factor. This gives an intuitive inverse DFT.

Definition 2 For a length N finite sequence $\{X[k]\}_0^{N-1}$ representing the DFT of a finite periodic signal $\{x[n]\}_0^{N-1}$, the inverse DFT is given by

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] e^{j \frac{2\pi}{N} kn}$$

Notice that the DFT and the IDFT are very similar in form. It turns out that the IDFT can be expressed as a DFT of $X^*[k]$. Namely

$$IDFT\{X[k]\} = \frac{1}{N} DFT\{X^*[k]\}^*$$

Further intuition for the DFT comes by relating it to the DTFT. Suppose we have a finite signal $x[n]$ which is 0 for $n < 0$ and $n > N - 1$. The DTFT of this signal is

$$X(\omega) = \sum_{n=-\infty}^{\infty} x[n] e^{-j\omega n} = \sum_{n=0}^{N-1} x[n] e^{-j\omega n}$$

Suppose we sample the DTFT at intervals of $\frac{2\pi}{N}k$, then the k th sample is given by

$$X[k] = X\left(\frac{2\pi}{N}k\right) = \sum_{n=0}^{N-1} x[n]e^{-j\frac{2\pi}{N}kn}$$

Thus we can think of the DFT as a N evenly spaced samples of the DTFT. One important point to notice is that while the DTFT is often centered around 0 (meaning it is plotted over a range from $-\pi$ to π), because we are summing from 0 to $N-1$ in the DFT, the DFT coefficients are centered around π , and thus they are plotted on a range fo $[0, 2\pi - \frac{2\pi}{N}]$

1.1 Convolution and the DFT

1.1.1 Circular Convolution

When the DFT coefficients of two signals are multiplied, the resulting coefficients describe a circular convolution of the original two signals.

$$x[n] \otimes y[n] \leftrightarrow X[k]Y[k]$$

Definition 3 A circular convolution between two finite sequences is given by

$$x[n] \otimes y[n] = \sum_{m=0}^{N-1} x[m]y[((n - m))_N]$$

The mechanics of the circular convolution are the same as that of the regular convolution, except the signal is circularly shifted as shown in fig. 1. A circular con-




Figure 1: A circular shift

volution is equivalent to a periodic convolution over a single period.

1.1.2 Linear Convolution with the DFT

Because multiplying DFT coefficients performs a specific case of convolution, we can compute a linear convolution using the circular convolution. Suppose we have two finite signals $\{x[n]\}_0^{L-1}$ and $\{h[n]\}_0^{P-1}$. The linear convolution of these two signals will be length $L + P - 1$, so in order to take an IDFT and get $L + P - 1$ samples, we need to take at least $N \leq L + P - 1$ points.

1. Pad each vector to length $L + P - 1$
2. Compute $X[k]H[k]$
3. Take the Inverse DFT

If N is smaller than $L + P - 1$, the result is akin to aliasing in the time domain. To see why, consider that the DFT coefficients are essentially the DTFS coefficients of the periodic extension of $x[n]$ (denote $\tilde{x}[n]$).

$$\tilde{x}[n] = \sum_{r=-\infty}^{\infty} x[n - rN]$$

If we compute the DTFT of each periodic extension, then

$$Y(e^{j\omega}) = X(e^{j\omega})H(e^{j\omega})$$

and the IDTFT of this will be

$$\tilde{y}[n] = \sum_{r=-\infty}^{\infty} y[n - rN].$$

Notice that if N is not large enough, then these copies will be overlapping (a.k.a aliasing). Since the DFT is just sampling the DTFT, the circular convolution will represent the true convolution so long as the copies don't overlap.

1.1.3 Block Convolutions

In a discrete time system, the input signal might have a very long length, making it impractical to be stored in a computer's memory or to compute the DFT of it all at once (especially if we have a real-time system). Thus to compute the output of a digital filter (with impulse response of length P), we need to compute the DFT in blocks shorter than the signal.

The first method of block convolution is the **overlap-add method**.

1. Decompose $x[n]$ into nonoverlapping segments of length L

$$x[n] = \sum_r x_r[n] \quad x_r[n] = \begin{cases} x[n] & rL \leq n \leq (r+1)L \\ 0 & \text{else.} \end{cases}$$

2. Since convolution is linear,

$$y[n] = x[n] * h[n] = \sum_r x_r[n] * h[n].$$

3. Zero pad $x_r[n]$ and $h[n]$ to length $N \geq L + P - 1$ to prevent time-domain aliasing
4. Compute the DFTs, multiply them, and take the inverse DFT.
5. The neighboring outputs overlap in the last $P - 1$ points, so add the overlapping sections together to get the final output

The other method of block convolution is the **overlap-save method**.

1. Divide $x[n]$ into sections of length L such that each section overlaps the previous by $P - 1$ points

$$x_r[n] = x[n + r(L - P + 1) - P + 1] \quad 0 \leq n \leq L - 1$$

2. Zero pad $x_r[n]$ and $h[n]$ to length $N \geq L + P - 1$ to prevent time domain aliasing.
3. Compute the DFTs, multiply the coefficients, and compute the inverse DFT.
4. The first $P - 1$ samples of the output will be incorrect, so we can discard them.

$$y[n] = \sum_{r=0}^{\infty} y_r[n - r(L - P + 1) + P - 1] \quad y_r[n] = \begin{cases} x_r[n] * h[n] & P - 1 \leq n \leq L - 1 \\ 0 & \text{else} \end{cases}$$

1.2 FFT

The DFT gives us an easy way to do convolutions. Unfortunately naively, computing it is an $O(N^2)$ operation because we must sum together N elements to compute N different coefficients. Thankfully, there is a fast algorithm which can compute the DFT in $O(N \log N)$ time so we can compute convolutions quickly.

Definition 4 *The fast fourier transform is an algorithm which computes the DFT efficiently in $O(N \log N)$ time.*

It works by exploiting properties of the N th roots of unity.

Definition 5 *The N th roots of unity are the complex roots of $W_N^N = 1$.*

$$W_N^k = e^{-j \frac{2\pi k}{N}}$$

The roots of unity have the following properties.

Theorem 1 *The N th roots of unity are conjugate symmetric.*

$$W_N^{N-kn} = W_N^{-kn} = (W_N^{kn})^*$$

Theorem 2 *The N th roots of unity are periodic in N .*

$$W_{kn} = W_N^{k(n+N)} = W_N^{(k+N)n}$$

Theorem 3 *When squared, the N th roots of unity become the $\frac{N}{2}$ th roots of unity.*

$$W_N^2 = W_{\frac{N}{2}}$$

Using theorems 1 to 3, we can take two approaches to the FFT: decimation in time, which splits $x[n]$ into smaller subsequences, and decimation in frequency which splits $X[k]$ into smaller subsequences.

1.2.1 Decimation in Time

The idea here is to break $x[n]$ into smaller subsequences. We assume that N is a power of 2 for simplicity.

$$X[k] = \sum_{n=0}^{N-1} x[n] W_N^{kn} = \sum_{n \text{ even}} x[n] W_N^{kn} + \sum_{n \text{ odd}} x[n] W_N^{kn}$$

We let $n = 2r$ and $n = 2r + 1$.

$$\begin{aligned} X[k] &= \sum_{r=0}^{\frac{N}{2}-1} x[2r] W_N^{2rk} + \sum_{r=0}^{\frac{N}{2}-1} x[2r+1] W_N^{k(2r+1)} \\ &= \sum_{r=0}^{\frac{N}{2}-1} x[2r] W_{\frac{N}{2}}^{rk} + W_N^k \sum_{r=0}^{\frac{N}{2}-1} x[2r+1] W_{\frac{N}{2}}^{kr} \end{aligned}$$

These are just the DFTs of the even and odd elements of the signal!

$$\therefore X[k] = G[k] + W_N^k H[k]$$

Both G and H are $\frac{N}{2}$ periodic, and notice that

$$W_N^{k+\frac{N}{2}} = e^{-j\frac{2\pi}{N}(k+\frac{N}{2})} = -W_N^k.$$

This means once we compute $G[k]$ and $H[k]$ we can compute $X[k]$ easily because

$$X[k] = G[k] + W_N^k H[k] \quad X\left[k + \frac{N}{2}\right] = G[k] - W_N^k H[k] \quad \text{for } k \in \left[0, \frac{N}{2}\right)$$

We can continue this relationship recursively downwards. Once we get too $N = 2$, we can represent this as a simple **butterfly operation**:

$$X[0] = x[0] + x[1] \quad X[1] = X[0] - X[1].$$

1.2.2 Decimation in Frequency

The decimation in frequency approach is very similar to the decimation in time approach except instead we split the frequency components

$$X[2r] = \sum_{n=0}^{\frac{N}{2}-1} x[n] W_N^{2rn} + \sum_{n=0}^{\frac{N}{2}-1} x\left[n + \frac{N}{2}\right] W_N^{2r(n+\frac{N}{2})} = W_N^{rn} \sum_{n=0}^{\frac{N}{2}-1} \left(x[n] + x\left[n + \frac{N}{2}\right] \right)$$

$$X[2r+1] = W_N^{rn} \sum_{n=0}^{\frac{N}{2}-1} \left(x[n] - x\left[n + \frac{N}{2}\right] \right)$$

2 Spectral Analysis




Figure 2: A Digital Signal Processing System

In real-world DSP systems, we are often converting a continuous time signal into a discrete one via sampling (as shown in fig. 2). Because input is constantly streaming in, we can't process all of it at once, especially for real-time applications. That is why we instead process blocks of length L at a time. This is accomplished by multiplying our sampled signal by a window function $w[n]$.

All window functions are real, even, and finite. This means they have real and symmetric DTFTs. The most simply window is a box window (a sinc in the frequency domain). When the signal is multiplied by a window, it amounts to a periodic convolution in the frequency domain.

Definition 6 A periodic convolution of the two spectra $X(e^{j\Omega})$ and $W(e^{j\Omega})$ is given by

$$V(e^{j\Omega}) = \frac{1}{2\pi} \int_{\langle 2\pi \rangle} X(e^{j\Omega}) W(e^{j(\Omega-\omega)}) d\omega$$

This periodic convolution means our choice of window function has an impact on our ability to resolve frequencies in the frequency domain.

1. If $W(e^{j\Omega})$ has a wide “main lobe” at the DC frequencies, then the spectrum of $V(e^{j\Omega})$ will be blurred
2. If $W(e^{j\Omega})$ has large “side lobes” at non DC frequencies, then **spectral leakage** occurs because larger frequencies start bleeding into lower ones.

Another factor which impacts our ability to resolve frequencies in frequency domain is the length of the window. Because an L point DFT samples the DTFT at L points, taking a larger window will resolve the DTFT better. If we don’t want to increase the window length (because doing so would increase the latency of our system), we can zero pad after windowing because zero padding has no impact on the DFT besides sampling the DTFT at more finely spaced samples.

2.1 Short Time Fourier Transform (STFT)

By looking at the DFT of a signal $x[n]$, we only get the frequency information across the entire duration of the signal. Likewise, just by looking at $x[n]$, we get no frequency information and only temporal information. The STFT is a tool to see both at once.

Definition 7 The short-time fourier transform localizes frequencies in time by creating a spectrogram, an image which shows what frequencies are occurring at what time.

$$X[n, \omega) = \sum_{m=-\infty}^{\infty} x[n+m] w[m] e^{j\omega m}$$

The result is discrete on the temporal axis and continuous on the frequency axis. Computing the STFT requires a window function $w[n]$.

Essentially, we slide a window function around $x[n]$ and compute the DTFT at every time point.

2.2 Discrete STFT

Definition 8 *The Discrete STFT is the discrete version of the STFT*

$$X[r, k] = \sum_{m=0}^{L-1} x[rR + m]w[m]W_N^{km}$$

Our window is of length L , R is how much we shift the window around, and $N \geq L$ is the DFT length we are taking.

Just like before, we take our window and slide it around the signal, computing DFTs at every time point. If $N > L$, then we are essentially computing a zero-padded DFT. The DSTFT produces a spectrogram which we can display digitally.

Definition 9 *The inverse DSTFT is given by*

$$x[rR + m]w_L[m] = \frac{1}{N} \sum_{k=0}^{N-1} X[n, k]W_N^{-km}$$

As long as the window is never 0 and the windows don't overlap,

$$x[n] = \frac{x[n - rL]}{w_L[n - rL]}.$$

2.3 Time-Frequency Uncertainty

When we compute the spectrogram of a signal, we can think of each coefficient as "tiling" the time-frequency plane. If we consider the normal N point DFT, each DFT coefficient is supported by N points in the time domain. Since the DFT samples the DTFT, it divides the range of $[0, 2\pi]$ into N segments of width $\frac{2\pi}{N}$. Each coefficient represents a section of this space, leading to a tiling looking like fig. 3 (for a 5 point DFT). Thinking about the DSTFT, each coefficient is computed using L points of the original signal. Each coefficient still represents intervals of $\frac{2\pi}{N}$ in the frequency axis, so it will lead to a tiling which looks like fig. 4. What these tilings show us is that because we have discretized time and frequency, there is some uncertainty regarding which times and frequencies each coefficient represents.

We can formalize this idea by considering a general transform. All transforms are really an inner product with a set of basis functions

$$T_x(\gamma) = \langle x(t), \phi_\gamma(t) \rangle = \int_{-\infty}^{\infty} x(t)\phi_\gamma^*(t)dt.$$




Figure 3: Time-Frequency tiling of a 5 point DFT




Figure 4: Time-Frequency tiling of the DSTFT

For each γ , $T_x(\gamma)$ is the projection of the signal onto the basis vector $\phi_\gamma(t)$. We can use Parseval's relationship to see that

$$\begin{aligned} T_x(\gamma) &= \langle x(t), \phi_\gamma(t) \rangle = \int_{-\infty}^{\infty} x(t) \phi_\gamma^*(t) dt \\ &= \frac{1}{2\pi} \int_{-\infty}^{\infty} X(j\Omega) \Phi_\gamma^*(j\Omega) d\Omega \\ &= \langle X(j\Omega), \frac{1}{2\pi} \Phi_\gamma(j\Omega) \rangle. \end{aligned}$$

This means that we can think of our transform not only as a projection of the signal onto a new basis, but also as a projection of the spectrum of our signal onto the spectrum of our basis function. Remember that projection essentially asks "How much of a signal can be explained by the basis". We can formalize this by looking

at the signal in a statistical sense and treat it as a probability distribution.

$$\begin{aligned} m_t &= \int_{-\infty}^{\infty} t |\psi(t)|^2 dt & m_{\Omega} &= \int_{-\infty}^{\infty} \Omega \frac{1}{2\pi} |\Psi(j\Omega)|^2 d\Omega \\ \sigma_t^2 &= \int_{-\infty}^{\infty} (t - m_t)^2 |\psi(t)|^2 dt & \sigma_{\Omega}^2 &= \int_{-\infty}^{\infty} (\Omega - m_{\Omega})^2 \frac{1}{2\pi} |\Psi(j\Omega)|^2 d\Omega \end{aligned}$$

m_t and m_{Ω} are the means of the signal and the spectrum. σ_t^2 and σ_{Ω}^2 are the variances. Together, they localize where our signal "lives" in the time-frequency spectrum. The uncertainty principle says

$$\sigma_t \sigma_w \geq \frac{1}{2}.$$

This means there is nothing we can do to get completely accurate time resolution and frequency resolution, and any decisions we make will lead to a tradeoff between them.

2.4 Wavelets

While the STFT gives us a better picture of a signal than a full-length DFT, one of its shortcomings is that each coefficient is supported by the same amount of time and frequency. Low frequencies don't change as much as high frequencies do, so a lower frequency needs to be resolved with more time support whereas a fast signal would requires less time support to resolve properly.

Definition 10 *The Wavelet transform finds coefficients which tile the time-frequency spectrum with different time and frequency supports using a mother and father wavelet.*

$$Wf(u, s) = \int_{-\infty}^{\infty} f(t) \frac{1}{\sqrt{s}} \Psi^* \left(\frac{t - u}{s} \right) dt$$

The Wavelet transform essentially makes all of the boxes in fig. 4 different sizes.

Definition 11 *The mother wavlet is a scaled bandpass filter $\Psi(t)$ used for the kernel of the wavelet transform. It must have the following properties:*

$$\int_{-\infty}^{\infty} |\Psi(t)|^2 dt = 1 \quad \int_{-\infty}^{\infty} \Psi(t) dt = 0$$

We need an infinite number of functions to fully represent all frequencies properly, but at a certain level, we don't care about our ability to resolve them better, so we stop scaling and use a low frequency function $\Phi(t)$ to "plug" the remaining bandwidth.

Definition 12 *The father wavelet is a low frequency function $\Phi(t)$ used to "plug" the remaining bandwidth not covered by the mother wavelet.*

2.4.1 Discrete Wavelet Transform

In discrete time, the wavelet transform becomes

$$d_{s,u} = \sum_{n=0}^{N-1} x[n] \Psi_{s,u}[n] \quad a_{s,u} = \sum_{n=0}^{N-1} x[n] \Phi_{s,u}[n]$$

The d coefficients are the detailed coefficients and are computed using the mother wavelet. They capture higher frequency information. The a coefficients are the approximate coefficients computed using the father wavelet. They represent lower frequency information. The time frequency tiling for the DWT looks like fig. 5. Notice how each wavelet coefficient is supported by a different amount of time




Figure 5: Time Frequency tiling of wavelets

and frequency. We can choose different mother and father wavelets to describe our signals depending on how we want to tile the time-frequency plane.

3 Sampling

3.1 Ideal Sampling

In order to work with continuous signals using a computer, we need to sample them. This means recording the value at particular points of time. During uniform sampling, we take samples at an even sampling period T_s so $x[n] = x_c(nT)$

(where x_c is our continuous signal). This is done by passing the signal through an Analog-To-Digital converter. From there we can do discrete time processing and reconstruct our signal by passing it through a Digital-to-Analog converter with reconstruction period T_r .




Figure 6: Uniform Sampling System

We mathematically model sampling as multiplication by an impulse train. Notice that if we were to take a signal $x(t)$ and multiply it by an impulse train, then we would get a series of impulses equal to $x(t)$ at the sampling points and 0 everywhere else. We can call this signal $x_p(t)$.

$$p(t) = \sum_{k=-\infty}^{\infty} \delta(t - kT)$$

$$x_p(t) = x(t)p(t) = \sum_{k=-\infty}^{\infty} x(t)\delta(t - kT)$$

In the Fourier Domain,

$$X_p(j\Omega) = \frac{1}{2\pi} X(j\Omega) * P(j\Omega)$$

$$P(j\Omega) = \frac{2\pi}{T} \sum_{k=-\infty}^{\infty} \delta(\Omega - k\Omega_s)$$

$$\therefore X_p(j\Omega) = \frac{1}{2\pi} \int_{-\infty}^{\infty} X(j\theta) P(j(\Omega - \theta)) d\theta = \frac{1}{T} \sum_{k=-\infty}^{\infty} X(j(\Omega - k\Omega_s))$$

What this tells us is that the Fourier Transform of our sampled signal is a series of copies of $X(j\Omega)$, each centered at $k\Omega_s$ where $\Omega_s = \frac{2\pi}{T}$. This is a good model because we can equivalently write the CTFT of the impulse train sampled signal as

$$X_p(j\Omega) = \int_{-\infty}^{\infty} \sum_{k=-\infty}^{\infty} x(t)\delta(t - kT) \sum_{k=-\infty}^{\infty} x(kT)e^{-jkT\Omega}.$$

Notice that this is just the DTFT of $x[n] = x(nT)$ if we set $\omega = \Omega T$.

$$X(e^{j\omega}) = \sum_{n=-\infty}^{\infty} x(nT)e^{-j\omega n} = X_p(j\Omega)|_{\Omega=\frac{\omega}{T}} = \frac{1}{T} \sum_{k=-\infty}^{\infty} X\left(\frac{\omega}{T} - k\frac{2\pi}{T_s}\right)$$

This means that the DTFT of our signal is just a bunch of shifted copies, and the frequency axis is scaled so $\Omega_s \rightarrow 2\pi$.

3.1.1 Nyquist Theorem

To analyze this further, we will stay in continuous time. Lets say that our original signal has the following Fourier Transform. Notice the signal is band-limited by Ω_M . There are two major cases: if $\Omega_s > 2\Omega_m$ and $\Omega_s < 2\Omega_m$.




Figure 7: Example of the spectrum of a bandlimited signal

Case One: $\Omega_s > 2\Omega_m$. As shown in fig. 8, the shifted copies of the original $X(j\Omega)$




Figure 8: When $\Omega_s > 2\Omega_m$

(shown in blue) do not overlap with each other or with the original copy. If we wanted to recover the original signal, we could simply apply a low pass filter to isolate the unshifted copy of $X(j\Omega)$ and then take the inverse Fourier Transform.

Case Two: $\Omega_s < 2\Omega_m$ Notice how in fig. 9, the shifted copies overlap with the




Figure 9: When $\Omega_s < 2\Omega_m$

original $X(\omega)$. This means in our sampled signal, the higher frequency information is bleeding in with the lower frequency information. This phenomenon is known as aliasing. When aliasing occurs, we cannot simply apply a low pass filter to isolate the unshifted copy of $X(\omega)$.

When $\Omega_s = 2\Omega_M$, then our ability to reconstruct the original signal depends on the shape of its Fourier Transform. As long as $X_p(jk\Omega_m)$ are equal to $X(j\Omega_m)$ and

$X(-j\Omega_m)$, then we can apply an LPF because we can isolate the original $X(j\Omega)$ and take its inverse Fourier Transform.

Remember that an ideal low pass filter is a square wave in the frequency domain and a sinc in the time domain. Thus if we allow

$$X_r(j\Omega) = X_p(j\Omega) \cdot \begin{cases} T & |\Omega| < \frac{\Omega_s}{2} \\ 0 & \text{else} \end{cases}$$

then our reconstructed signal will be

$$x_r(t) = x_p(t) * \text{sinc}\left(\frac{t}{T}\right) = \sum_{n=-\infty}^{\infty} X(nT) \text{sinc}\left(\frac{t - nT}{T}\right).$$

This is why we call reconstructing a signal from its samples "sinc interpolation." This leads us to formulate the Nyquist Theorem.

Theorem 4 (Nyquist Theorem) *Suppose a continuous signal x is bandlimited and we sample it at a rate of $\Omega_s > 2\Omega_m$, then the signal $x_r(t)$ reconstructed by sinc interpolation is exactly $x(t)$*

3.1.2 Discrete Time Processing of a Continuous Time Signal

As long as the DT system we apply is LTI, the overall CT system will be linear too, but it will not necessarily be time invariant because sampling inherently depends on the signal's timing. If we want to find the overall CT transfer function ($\omega = \Omega T$) of a system like that depicted in fig. 6.

$$\begin{aligned} Y_d(e^{j\omega}) &= H_d(e^{j\omega})X_d(e^{j\omega}) = H_d(e^{j\omega})X_p\left(\frac{\omega}{T}\right) \\ Y_p(j\Omega) &= Y_d(e^{j\Omega T}) = H_d(e^{j\Omega T})X_p(j\Omega) \\ Y(j\Omega) &= \begin{cases} T & |\Omega| < \frac{\Omega_s}{2} \\ 0 & |\Omega| \geq \frac{\Omega_s}{2} \end{cases} \cdot Y_p(j\Omega) = \begin{cases} TH_d(e^{j\Omega T})X_p(j\Omega) & |\Omega| < \frac{\Omega_s}{2} \\ 0 & |\Omega| \geq \frac{\Omega_s}{2} \end{cases} \end{aligned}$$

Assuming that the Nyquist criteria is met holds,

$$\begin{aligned} X_p(j\Omega) &= \frac{1}{T}X(j\Omega) \\ \therefore Y(j\Omega) &= \begin{cases} H_d(e^{j\Omega T})X(j\Omega) & |\Omega| < \frac{\Omega_s}{2} \\ 0 & |\Omega| \geq \frac{\Omega_s}{2} \end{cases} \\ \therefore H_{system} &= \begin{cases} H_d(e^{j\omega T}) & |\Omega| < \frac{\Omega_s}{2} \\ 0 & |\Omega| \geq \frac{\Omega_s}{2} \end{cases} \end{aligned}$$

This shows us that as long as the Nyquist theorem holds, we can process continuous signals with a discrete time LTI system and still have the result be LTI.

3.1.3 Continuous Time Processing of Discrete Time Signals

While not useful in practice, it can be useful to model a discrete time transfer function in terms of Continuous Time processing (e.g a half sample delay).




Figure 10: Continuous Time processing of a Discrete Time signal

Similar to the analysis of DT processing of a CT signal, we can write the discrete transfer function in terms of the continuous function. Our continuous signal will be bandlimited after reconstruction.

$$X(j\Omega) = \begin{cases} TX_d(e^{j\omega})|_{\omega=\Omega T} & |\Omega| \leq \frac{\Omega_s}{2} \\ 0 & \text{otherwise} \end{cases}$$

This means our reconstructed signal $Y(j\Omega) = H(j\Omega)X(j\Omega)$ is also bandlimited, so we can say that

$$Y_d(e^{j\Omega}) = H(j\Omega)|_{\Omega=\frac{\omega}{T}} X(e^{j\omega})$$

3.1.4 Downsampling

When we downsample a signal by a factor of M , we create a new signal $y[n] = x[nM]$ by taking every M th sample. What this means conceptually is that we are reconstructing the continuous signal and then sampling it at a slower rate MT where T was the original sampling rate. If x_c is the original continuous time signal and x_d is the sampled signal, then the downsampled signal $y[n]$ will be

$$y[n] = x[nM] = x_c(nMT) \implies Y(e^{j\omega}) = \frac{1}{MT} \sum_{k=-\infty}^{\infty} X_c\left(\frac{\omega}{NT} - k\frac{2\pi}{NT}\right).$$

If we re-index and let $k = Mp + m$ for $m \in [0, N - 1], p \in \mathbb{Z}$

$$Y(e^{j\omega}) = \frac{1}{M} \sum_{m=0}^{M-1} X_d(e^{j\frac{\omega - 2\pi m}{M}}).$$

What this means is to obtain the new DTFT, we need to scale the frequency axis so $\frac{\pi}{M} \rightarrow \pi$. To prevent aliasing when this happens, we include an LPF before the downsample step.




Figure 11: Downsampling

3.1.5 Upsampling

When we upsample a signal by a factor of L , we are interpolating between samples. Conceptually, this means we are reconstructing the original continuous time signal and resampling it at a faster rate than before. First we place zeros in between samples, effectively expanding our signal.

$$x_e[n] = \begin{cases} x\left[\frac{n}{L}\right] & n = 0, \pm L, \pm 2L, \dots \\ 0 & \text{otherwise} \end{cases}$$

$$X_e(e^{j\omega}) = \sum_{n=-\infty}^{\infty} x_e[n] e^{-j\omega n} = \sum_{m=-\infty}^{\infty} x[m] e^{-j\omega mL} = X(e^{j\omega L})$$

Then we interpolate by convolving with a sinc.

$$y[n] = x_e[n] * \text{sinc}\left(\frac{n}{L}\right) = \sum_{k=-\infty}^{\infty} x[k] \text{sinc}\left(\frac{n - kL}{L}\right)$$

In the frequency domain, this looks like compressing the frequency axis so $\pi \rightarrow \frac{\pi}{L}$ and then taking a low pass filter. The gain of L is used to scale the spectrum so it




Figure 12: Upsampling operation

is identical to if we had sampled the continuous signal at a rate of $\frac{T}{L}$.

3.2 Multi-Rate Signal Processing

In order to resample a signal to a rate $T' = \frac{MT}{L}$ where T is the original sampling rate, we can do this by upsampling then downsampling our signal. Notice that we only need one LPF to take care of both anti-aliasing and interpolation.




Figure 13: Resampling

3.2.1 Exchanging Filter Order During Resampling

Notice that resampling with a very small change wastes a lot of computation. For example, resampling with $T' = 1.01T$ would upsample by 100 and then throw away most of those samples when we downsample. Thus it would be useful to exchange the order of operations when resampling to save computation. During




Figure 14: Interchanging an upsampling operation

upsampling, we convolve our filter with a bunch of zeros caused by the expansion. Convolution with 0's is unnecessary, so instead we could convolve with a compressed version of the filter. Notice the results will be the same as long as $H(z^{\frac{1}{L}})$ is a rational function. During downsampling, we do a convolution and then throw




Figure 15: Interchanging a downsampling operation

away most of our results. It would be much more efficient to instead compute only the quantities we need. This is accomplished by downsampling first and then convolving. Just like before, the results are only going to be the same if $H(z^{\frac{1}{M}})$ is a rational function.

3.2.2 Polyphase Decomposition

The problem with interchanging filters is that it is not always possible. Most filters are not compressible. However, we can get around this issue and still get the efficiency gains of interchanging filter orders by taking a polyphase decomposition of our filters. First notice that $h[n]$ can be written as a sum of compressible filters.

$$h[n] = \sum_{k=0}^{M-1} h_k[n - k]$$

This means if we let $e_k[n] = h_k[nM]$, we can utilize the linearity of convolution to build a bank of filters. Now each of our filters is compressible, so we can switch the




Figure 16: Example of decomposing a filter: $M=2$




Figure 17: A filter bank




Figure 18: Filter bank but with the downsampling done first

order of downsampling and filtering while maintaining the same output. Now for any filter, we can compute only what we need, so the result is correct and efficiently obtained.

3.3 Practical Sampling (ADC)

Unfortunately, ideal analog to digital conversion is not possible for a variety of reasons. The first is that not all signals are bandlimited (or there may be noise outside of the bandwidth). Moreover, computers only have finite precision, so we cannot represent the full range of values that a continuous signal might take on with a finite number of bits per sample. The solution to the first issue is to include a “anti-aliasing” filter before the sampler. The solution to the second issue is to

quantize. However, sharp analog filters are difficult to implement in practice. To




Figure 19: Sampling with quantization

deal with this, we could make the anti-aliasing filter wider, but this would add noise and interference. If we keep the cutoff frequency the same, then we could alter part of the signal because our filter is not ideal. A better solution is to do the processing in Discrete Time because we have more control. We also sample higher than the Nyquist Rate and then downsample it to the required rate.




Figure 20: A practical sampling system with quantization and anti-aliasing

3.3.1 Quantization

If we have a dynamic range of X_m (i.e $2X_m$ is the length of the range of values we can represent), then our step between quantized values is $\Delta = \frac{X_m}{2^B}$, assuming we are representing our data as 2's complement numbers with B bits. We model the error caused by quantization as additive noise. Our quantized signal $\hat{x}[n]$ is described by

$$\hat{x}[n] = x[n] + e[n] \quad -\frac{\Delta}{2} \leq e[n] \leq \frac{\Delta}{2}$$

We do this under the following assumptions:

1. $e[n]$ is produced by a stationary random process
2. $e[n]$ is not correlated with $x[n]$
3. $e[n]$ is white noise ($e[n]$ is not correlated with $e[m]$)
4. $e[n] \sim U \left[-\frac{\Delta}{2}, \frac{\Delta}{2} \right]$

For rapidly changing signals with small Δ , this assumptions hold, and they are useful in modeling quantization error. Since $\Delta = 2^{-B} X_m$

$$\sigma_e^2 = \frac{\Delta^2}{12} = \frac{2^{-2B} X_m^2}{12}$$

This means our Signal to Noise Ratio for quantization is

$$SNR_Q = 10 \log \left(\frac{\sigma_x^2}{\sigma_e^2} \right) = 6.02B + 10.8 - 20 \log \left(\frac{X_m}{\sigma_s} \right)$$

What this tells us is that every new bit we add gives us 6dB in improvement. It also tells us that we need to adapt the range of quantization to the RMS amplitude of the signal. This means there is a tradeoff between clipping and quantization noise. When we oversampling our signal, we can further limit the effects of quantization noise because this noise will be spread out over more frequencies and the LPF will eliminate noise outside the signal bandwidth. This makes $\frac{\sigma_e^2}{M}$ the new noise variance (if we oversample by M). Thus we can modify the SNR_Q equation

$$SNR_Q = 6.02B + 10.8 - 20 \log \left(\frac{X_m}{\sigma_s} \right) + 10 \log M.$$

This shows that doubling M yields a 3dB improvement (equivalent to 0.5 more bits).

3.4 Practical Reconstruction (DAC)

In the ideal case, we reconstruct signals by converting them to impulses and then convolving with a sinc. However, impulses are require lots of power to generate, and sines are infinitely long, so it is impractical to design an analog system to do this. Instead, we use an interpolation like Zero-Order-Hold to pulses and then filter with a reconstruction filter.




Figure 21: Practical Reconstruction

$$X_r(j\Omega) = \overbrace{H_r(j\Omega) T e^{-j\Omega \frac{T}{2}} \text{sinc} \left(\frac{\Omega}{\Omega_s} \right)}^{\text{Zero Order Hold}} \overbrace{\frac{1}{T} \sum_{k=-\infty}^{\infty} X(j(\Omega - k\Omega_s))}^{\text{Sampled Signal}}$$

We design $H_r(j\Omega)$ such that $H_r(j\omega)H_0(j\omega)$ is approximately an LPF.

4 Filtering

4.1 Transform Analysis of LTI Systems

LTI filters are characterized by their impulse response. The two broad categories of LTI systems are those with finite impulse responses (FIR) and those with infinite impulse responses (IIR). LTI systems are frequently characterized by linear constant-coefficient difference equations which look as follows:

$$\sum_{k=0}^N a_k y[n-k] = \sum_{k=0}^M x[n-k].$$

Definition 13 The system function $H(z)$ is the z -transform of the impulse response of the system. For LCCDE's, it is a ratio of polynomials in z^{-1} .

$$H(z) = \frac{\sum_{k=0}^M b_k z^{-k}}{\sum_{k=0}^N a_k z^{-k}} = \frac{b_0 \prod_{k=1}^M (1 - c_k z^{-1})}{a_0 \prod_{k=1}^N (1 - d_k z^{-1})}$$

We call c_k (the roots of the numerator) the **zeros** of the system and d_k (the roots of the denominator) the **poles** of the system.

Definition 14 The Magnitude Response $|H(e^{j\omega})|$ describes how the system will scale a complex exponential.

$$|H(e^{j\omega})| = \frac{|b_0| \prod_{k=1}^M |1 - c_k e^{-j\omega}|}{|a_0| \prod_{k=1}^N |1 - d_k e^{-j\omega}|}$$

Definition 15 The Phase Response $\arg[H(e^{j\omega})]$ describes how the system will shift the phase of a complex exponential.

$$\arg[H(e^{j\omega})] = \arg[b_0] + \sum_{k=1}^M \arg[1 - c_k e^{-j\omega}] - \arg[a_0] - \sum_{k=1}^N \arg[1 - d_k e^{-j\omega}]$$

Definition 16 The Group Delay of the system $\text{grad}[H(e^{j\omega})]$ tells us how much a complex exponential will be delayed.

$$\text{grad}[H(e^{j\omega})] = -\frac{d}{d\omega} \arg[H(e^{j\omega})] = \sum_{k=1}^M \text{grad}[1 - c_k e^{-j\omega}] - \sum_{k=1}^N \text{grad}[1 - d_k e^{-j\omega}]$$

We can systematically analyze these by drawing a vector from $e^{-j\omega}$ to each d_k or c_k and analyze each one individually. For example, if we look at one pole in the magnitude response

$$|1 - d_k e^{-j\omega}| = |e^{j\omega} - d_k| = |v_k|.$$

In general, the effects of poles and zeros on each of these quantities is described by the following table.

	Magnitude Response	Phase Response	Group Delay
Poles	Increase	Phase Lag	Increase
Zeros	Decrease	Phase Advance	Decrease

These effects are larger when c_k or d_k are close the unit circle (i.e $|c_k|, |d_k| \approx 1$).

4.2 All Pass Systems

Definition 17 All pass systems are those where $|H(e^{j\omega})| = k$ where k is some constant gain.

$$H(z) = k \prod_{k=1}^{M_r} \frac{z^{-1} - d_k}{1 - d_k z^{-1}} \prod_{k=1}^{M_c} \frac{(z^{-1} - e_k^*)(z^{-1} - e_k)}{(1 - e_k z^{-1})(1 - e_k^* z^{-1})}$$

Their Z-transform has the real poles d_k cancelled by real zeros, and complex poles e_k cancelled by the conjugates e_k^* .

Theorem 5 If an All-Pass system is stable, then

1. $\text{grad}[H(e^{j\omega})] > 0 \implies \text{Causal}$
2. $\arg[H(e^{j\omega})] \leq 0 \implies \text{Phase Lag}$

4.3 Minimum Phase Systems

Definition 18 A stable and causal system $H(z)$ whose inverse $\frac{1}{H(z)}$ is also stable and causal is called a Minimum Phase System.

What this means is that all poles and zeros must be inside the unit circle, and the region of convergence is right sided. Minimum phase systems are called minimum phase because of all $H(z)$ with the same magnitude response, a minimum phase system has the minimum phase and the minimum group delay.

Theorem 6 Any stable and causal system can be decomposed into a minimum phase system and an all-pass system.

$$H(z) = H_{min}(z)H_{ap}(z)$$

This is useful because if a signal undergoes a distortion, we can at least undo the minimum phase part of it (since H_{min} has a guaranteed inverse).

4.4 Generalized Linear Phase Systems

Definition 19 A linear phase system is one with constant group delay.

$$H(e^{j\omega}) = A(e^{j\omega})e^{-j\alpha\omega} \implies \text{grd}[H(e^{j\omega})] = \alpha$$

Note that $A(e^{j\omega})$ is a real function.

Definition 20 A generalized linear phase system has frequency response given by

$$H(e^{j\omega}) = A(e^{j\omega})e^{-j\alpha\omega+\beta} \implies \text{grd}[H(e^{j\omega})] = \alpha$$

If we limit ourselves to using FIR filters, then a GLP system must have either even or odd symmetry, meaning for some M

$$h[n] = h[M - n] \text{ or } h[n] = -h[M - n].$$

This restricts us to 4 different filter types.

	Symmetry	M	Filter Types	Notes
Type I	Even	Even	All	
Type II	Even	Odd	Low Pass	$H(e^{j\pi}) = 0$
Type III	Odd	Even	Bandpass	$H(e^{j0}) = H(e^{j\pi}) = 0$
Type IV	Odd	Odd	High	$H(e^{j0}) = 0$

Because of their symmetry, FIR systems are limited in where their zeros.

$$\text{Type I, II: } H(z) = z^{-M}H(z^{-1}) \quad \text{Type III, IV: } H(z) = -z^{-M}H(z^{-1})$$

In other words, if $a = re^{j\theta}$ is a zero, then $\frac{1}{a^*}$ is too. We can decompose GLP systems into a minimum phase, maximum phase, and unit circle system.

4.5 Filter Design

The idea of filter design is to take a desired frequency response and design a filter which has that frequency response. Some frequency responses can only be described by IIR systems which are impractical for real applications, so we make various tradeoffs when we design FIR filters to implement in our systems. We also like our filters to be causal because it makes them usable in real-time systems.

A M th order causal filter has $M + 1$ coefficients.

Definition 21 *The time-bandwidth product describes how sinc-like a filter looks like.*

$$TBW = (M + 1) \frac{\omega_c}{\pi}$$

The TBW is also the number of zero-crossings in the impulse response (including the zeros crossings at the end of the filter). To generate a High Pass filter, we can design a Low-Pass filter and then modulate it

$$h_{hp}[n] = (-1)^n h_{lp}[n].$$

We can do the same for a bandpass filter

$$h_{bp}[n] = 2h_{lp} \cos(\omega_0 n).$$

4.5.1 Windowing Method

One way to generate a filter which matches a desired frequency response is through windowing.

1. Choose a desired frequency response (often non-causal and IIR)
2. Window the Impulse Response
3. Module to shift the impulse response to make it casual

The length of the window impacts the transition width (how long it takes to transition). A longer window means a smaller width. The window type will impact the ripples in the frequency response. The choice of of window and its sidelobes impact these magnitudes.

4.5.2 Optimal Filter Design

With optimal filter design, we set up constraints to find a $H_d(e^{j\omega})$ which approximates $H(e^{j\omega})$ based on our optimization requirements. In general, we have some regions $W_c \subseteq [0, \pi]$ that we care about and other regions that we don't care about. We can first design a noncausal filter $H(\tilde{e}^{j\omega})$ and then shift it to make it causal. We do this by sampling and discretizing the frequency response to $\omega_k = k\frac{\pi}{P}$ where $-\pi \leq \omega_1 \leq \dots \leq \omega_p \leq \pi$. We choose P to be sufficiently big and make sure the $\omega_k \in \omega_c$ (the region we care about). In a least squares setup, we can solve

$$\underset{\tilde{h}}{\operatorname{argmin}} ||A\tilde{h} - b||^2 \quad A = \begin{bmatrix} e^{-j\omega_1(-N)} & \dots & e^{-j\omega_1(N)} \\ \vdots & \ddots & \vdots \\ e^{-j\omega_p(-N)} & \dots & e^{-j\omega_p(N)} \end{bmatrix} \quad \vec{b} = \begin{bmatrix} H(e^{j\omega_1}) \\ \vdots \\ H(e^{j\omega_p}) \end{bmatrix}$$

Other possible optimizations are Weighted Least Squares or Chebyshev Design.

$$\textbf{WLS:} \min \int_{-\pi}^{\pi} W(\omega) |H(e^{j\omega}) - H_d(e^{j\omega})|^2 d\omega$$

$$\textbf{Chebyshev:} \min_{w \in W_c} \max |H(e^{j\omega} - H_d(e^{j\omega}))|$$

Another optimization technique is the min-max ripple design where we try and control the deviations of the filter from the desired response. We can set up a linear program to do this for us. For example, if we were designing a low pass filter, we could write the LP

$$\begin{aligned} & \min \delta \\ & 1 - \delta \leq H_d(e^{j\omega_k}) \leq 1 + \delta \quad 0 \leq \omega_k \leq \omega_c \\ & -\delta \leq H_d(e^{j\omega_k}) \leq \delta \quad \omega_c \leq \omega_k \leq \pi \\ & \delta > 0 \end{aligned}$$